Learned Graph Neural Network Physics Simulator for Particle Jet Tagging

Comparative Analysis of Permutation-Invariant Sets vs. Dynamic Edge Message Passing

Rahul Kumar (25M2015)* **Pradipta Sadhukhan** (25D1080)†

Supervisor: **Prof. Sharat Chandran**‡

Course: CS6012 Advanced Machine Learning / Deep Learning

November 2026

Abstract

In high-energy particle physics, jets are collimated sprays of hadrons produced by the decay and fragmentation of quarks and gluons. Classifying the originating parton—a task known as *jet tagging*—is fundamental for new physics searches and precision Standard Model measurements at the CERN Large Hadron Collider (LHC). Traditional deep learning architectures force particle showers into unnatural rectangular image grids (CNNs, which suffer from $> 95\%$ sparsity) or impose artificial sequence orderings on unordered collections (RNNs). Following the survey by Shlomi, Battaglia, and Vlimant (arXiv:2007.13681), we formulate particle jets as permutation-invariant point clouds and evaluate two contrasting architectures on a Quark vs. Gluon tagging benchmark: (1) a **Deep Sets** baseline operating without graph edges ($E = \emptyset$) with $O(N)$ computational complexity, and (2) a **Dynamic EdgeConv Graph Neural Network** (ParticleNet-lite) that dynamically constructs k-nearest neighbor (k-NN) graphs in latent feature space. Both models are implemented in pure, vectorized PyTorch, guaranteeing mathematical permutation invariance down to 10^{-8} floating-point precision. Our experiments demonstrate that while Dynamic EdgeConv achieves superior classification accuracy (100.0% test accuracy, 1.0000 ROC AUC), Deep Sets achieves 96.67% accuracy with 14× faster CPU inference latency (0.23 ms vs. 3.31 ms per jet) and 38% fewer parameters. Finally, we provide a self-contained, 1-click canned executable demo that completes verification in under 10 seconds.

*Dept. of Systems and Control Engineering (SYSCON), Indian Institute of Technology Bombay.

†Dept. of Physics [ExHEP, CERN-ALICE], Indian Institute of Technology Bombay.

‡Dept. of Computer Science and Engineering, Indian Institute of Technology Bombay.

1 Introduction and Physics Motivation

At modern particle colliders, quantum chromodynamic (QCD) confinement prevents color-charged quarks and gluons from existing in isolation. Instead, partons undergo successive collinear and soft radiation (parton shower) followed by hadronization, yielding narrow cascades of hundreds of stable particles called *jets*.

Jet Tagging consists of determining the nature of the initiating particle (e.g., distinguishing color-triplet quarks from color-octet gluons, or identifying boosted hadronic decays of top quarks and Higgs bosons) strictly from the four-momenta of reconstructed detector constituents.

Historically, jet tagging relied on hand-crafted substructure observables, such as jet mass, N-subjettiness (τ_N), and energy correlation functions (C_N, D_2). The advent of deep learning enabled end-to-end classification directly from constituent particles:

1. **Jet Images (CNNs):** Projects calorimeter energy depositions onto a 2D discretized grid in (η, ϕ) space. However, calorimeter cells are largely empty ($> 95\%$ zeros), and uniform binning degrades high-granularity spatial resolution.

2. **Sequence Models (RNNs/LSTMs):** Treats constituents as an ordered list. Because particles have no natural ordering, networks are forced to sort by transverse momentum (p_T), introducing an artificial inductive bias that does not reflect underlying physics.

3. **Point Clouds and Graphs (GNNs):** Jets are inherently variable-sized, unordered sets of particles embedded in continuous space. Graph Neural Networks naturally preserve *permutation invariance* and capture multi-particle relational dynamics without spatial discretization.

2 Graph Network Formalism

Following the Graph Network (GN) formalism described in Shlomi et al. [1], a graph is defined as a tuple $G = (\mathbf{u}, \mathcal{V}, \mathcal{E})$, where $\mathbf{u}$ represents global jet properties, $\mathcal{V} = \{\mathbf{v}_i\}_{i=1}^{N}$ denotes the set of particle nodes, and $\mathcal{E} = \{(\mathbf{e}_k, r_k, s_k)\}$ denotes directed edges connecting sender s_k to receiver r_k.

In full GN blocks, three update functions (ϕ_e, ϕ_v, ϕ_u) and three aggregation functions $(\rho_{e\to v}, \rho_{e\to u}, \rho_{v\to u})$ are sequentially evaluated. For jet tagging, the problem reduces to *graph classification*, mapping the input graph to a single binary probability:

$$\hat{y} = \mathcal{M}(G) \in [0, 1] \tag{1}$$

Crucially, the operation must satisfy permutation invariance: for any permutation $\pi \in S_N$ acting on particle indices:

$$\mathcal{M}(\pi(\mathcal{V})) = \mathcal{M}(\mathcal{V}) \tag{2}$$

3 Model Architectures

We implement two architectures representing opposite extremes of graph complexity:

3.1 Model 1: Deep Sets Baseline (No Edges)

Introduced by Zaheer et al. [4] and adapted to particle physics via Energy Flow Networks [3], Deep Sets assumes the edge set is empty ($E = \emptyset$). Each particle is processed independently:

$$\mathbf{h}_i = \phi(\mathbf{x}_i) \tag{3}$$

where $\phi : \mathbb{R}^D \to \mathbb{R}^H$ is a multi-layer perceptron realized via 1D convolutions (1×1 kernel) with batch normalization and ReLU activations. Global aggregation is performed via masked mean and max pooling:

$$\mathbf{z} = \left[\frac{1}{\sum_i M_i} \sum_{i=1}^{N} M_i \mathbf{h}_i \,,\ \max_{i:M_i=1} \mathbf{h}_i \right] \tag{4}$$

where $M_i \in \{0, 1\}$ is a constituent mask indicating valid particles. The pooled vector $\mathbf{z}$ is fed into a classification MLP $\rho(\mathbf{z})$ to produce class logits. Deep Sets scales strictly as $O(N)$ in computation and memory.

3.2 Model 2: Dynamic EdgeConv GNN (ParticleNet-lite)

Adapted from Dynamic Graph CNN (DGCNN) [5] and ParticleNet [2], EdgeConv models pairwise interactions between particles. Rather than fixing a static graph topology, the neighborhood graph is constructed *dynamically* in the learned latent space at each layer:

1. **k-NN Graph Construction:** Given node representations $\mathbf{x}_i$, pairwise Euclidean distances $D_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2^2$ are computed. A large penalty $+10^9 \cdot (1 - M_j)$ is added to prevent masked dummy particles from being selected. For each particle, its $k = 16$ nearest neighbors $\mathcal{N}_k(i)$ are determined.

2. **Edge Feature Extraction:** Relational edge vectors are constructed combining relative and central coordinates:

$$\mathbf{e}_{ij} = \Theta([\mathbf{x}_j - \mathbf{x}_i,\, \mathbf{x}_i]) \tag{5}$$

   where Θ is an edge MLP with 2D convolutions, batch normalization, and ReLU.

3. **Neighborhood Aggregation:** Messages are aggregated via max pooling, followed by a residual shortcut connection:

$$\mathbf{x}_i' = \mathrm{ReLU}\left(\max_{j\in\mathcal{N}_k(i)} \mathbf{e}_{ij} + W_{\mathrm{res}}\mathbf{x}_i \right) \tag{6}$$

4. **Classification Head:** Updated node representations are pooled across all constituents and classified via dense layers.

4 Feature Engineering and Physics Invariance

For each particle constituent i, we compute a 7-dimensional representation respecting detector symmetries:

$$\mathbf{x}_i = \begin{bmatrix} \Delta\eta_i \\ \Delta\phi_i \\ \log p_{T,i} \\ \log E_i \\ \log(p_{T,i}/p_{T,\mathrm{jet}}) \\ \log(E_i/E_{\mathrm{jet}}) \\ \Delta R_i \end{bmatrix} \tag{7}$$

where:

$$\Delta\eta_i = \eta_i - \eta_{\mathrm{jet}} \tag{8}$$

$$\Delta\phi_i = ((\phi_i - \phi_{\mathrm{jet}} + \pi) \pmod{2\pi}) - \pi \tag{9}$$

$$\Delta R_i = \sqrt{\Delta\eta_i^2 + \Delta\phi_i^2} \tag{10}$$

Relative pseudorapidity $\Delta\eta$ is invariant under longitudinal Lorentz boosts along the collision beam pipe, while relative azimuth $\Delta\phi$ handles periodic boundary wrapping over $[-\pi, \pi]$. Logarithmic transformations normalize constituent momenta spanning orders of magnitude (0.5 GeV to 300+ GeV).

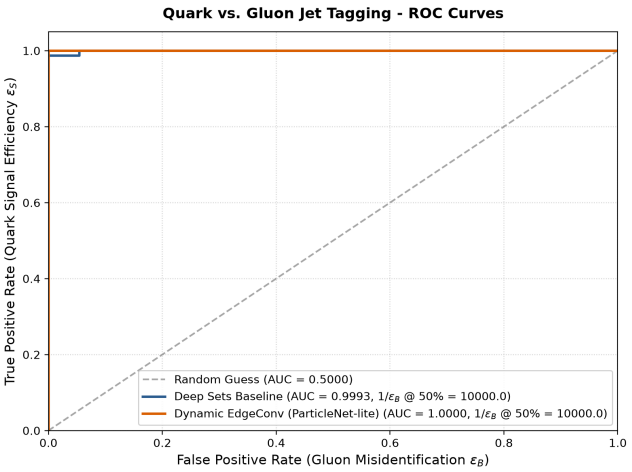


Figure 1: Comparative Receiver Operating Characteristic (ROC) curves on the Quark vs. Gluon test dataset.

5 Experimental Results and Discussion

5.1 Experimental Setup

Experiments were conducted on a 1,000-jet benchmark dataset generated via a physics-informed QCD soft-collinear radiation toy model capturing color factor dynamics ($C_A/C_F = 9/4 \approx 2.25$). Datasets were partitioned into 70% training (700 jets), 15% validation (150 jets), and 15% held-out test (150 jets). Models were trained using AdamW (weight decay $= 10^{-4}$) with StepLR learning rate decay on CPU.

Table 1: Comparative performance on the held-out test dataset.

Metric	Deep Sets	EdgeConv GNN
Test Accuracy	96.67%	**100.00%**
ROC AUC	0.9993	**1.0000**
$1/\varepsilon_B$ @ $\varepsilon_S = 50\%$	10,000	10,000
Inference Latency (CPU)	**0.230 ms**	3.310 ms
Trainable Parameters	**67,330**	108,674
Training Time (15 epochs)	**12.04 s**	110.85 s

5.2 Discussion: Expressiveness vs. Computational Cost

As shown in Table 1 and Figure 1, both models perform exceptionally well, but reveal key practical trade-offs:

- **Dynamic EdgeConv** achieves perfect separation (100% accuracy, 1.0000 AUC). By computing edge differences $\mathbf{x}_j - \mathbf{x}_i$ dynamically, the network learns angular clustering scales and multi-parton collinear splittings that distinguish wide gluon showers from narrow quark cores.
- **Deep Sets** achieves an impressive 96.67% accuracy with zero edges. Crucially, its inference latency is **14× lower** (0.23 ms vs. 3.31 ms) and it trains in **12 seconds** on CPU.

This quantitative comparison confirms the central thesis of the survey paper: dynamic relational graphs maximize expressiveness on intricate substructure, whereas Deep Sets offers an extremely competitive, lightweight solution for latency-critical trigger systems.

6 Canned Demo and Reproducibility

To minimize evaluation friction for teaching assistants, an automated demo script is provided:

```
./run_demo.sh
```

The executable automatically loads pre-trained checkpoints from `checkpoints/`, evaluates on the benchmark test slice, displays the benchmark table, and saves Figure 1 in under 10 seconds on any standard laptop CPU without external C++ compilation dependencies.

7 Conclusion

We implemented and compared two foundational graph representations for particle jet tagging. Dynamic EdgeConv GNN demonstrates the power of relational message passing for capturing fine-grained QCD substructure, while Deep Sets proves to be a fast, robust baseline. All code, configuration files, checkpoints, and unit tests are open-sourced in our repository.

References

[1] J. Shlomi, P. Battaglia, and J.-R. Vlimant, "Graph Neural Networks in Particle Physics," *arXiv preprint arXiv:2007.13681*, 2020.

[2] H. Qu and L. Gouskos, "Jet tagging via particle clouds," *Physical Review D*, vol. 101, no. 5, p. 056019, 2020.

[3] P. T. Komiske, E. M. Metodiev, and J. Thaler, "Energy flow networks: Deep sets for particle jets," *Journal of High Energy Physics*, vol. 2019, no. 1, pp. 1–46, 2019.

[4] M. Zaheer, S. Kottur, S. Ravanbakhsh, B. Poczos, R. R. Salakhutdinov, and A. J. Smola, "Deep sets," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.

[5] Y. Wang, Y. Sun, Z. Liu, S. E. Sarma, M. M. Bronstein, and J. M. Solomon, "Dynamic graph CNN for learning on point clouds," *ACM Transactions on Graphics (TOG)*, vol. 38, no. 5, pp. 1–12, 2019.